

Section 9 Introduction to Data Science


What is Data Science?

Three correlated concepts:

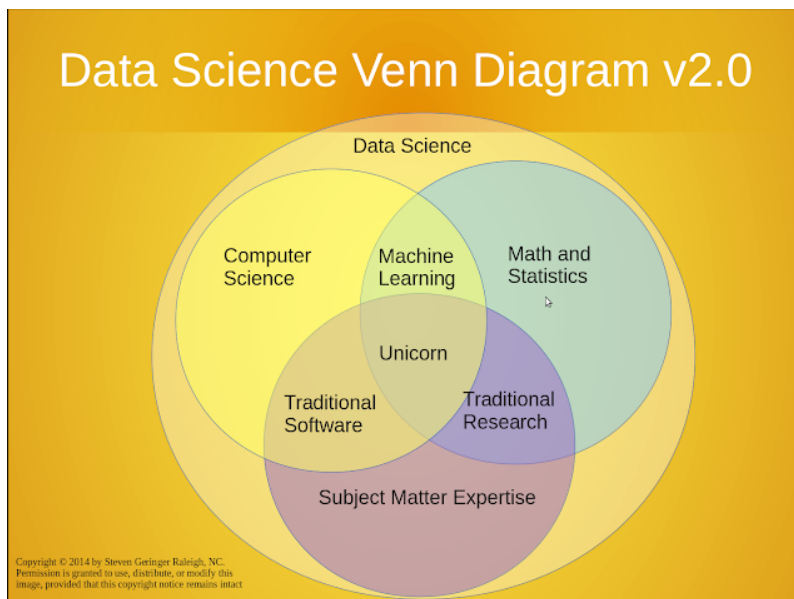
- Data Science
- Artificial Intelligence
- Machine Learning

Battle of the Data Science Venn Diagrams


The original Venn diagram from Drew Conway:





Another diagram from Steven Geringer:



Another version:



Perhaps the reality should be:



David Robinson's Auto-pilot example:

- machine learning: **predict** whether there is a stop sign in the camera

- artificial intelligence: design the **action** of applying brakes (either by rules or from data)
- data science: provide the **insights** why the system does not work well after sunrise

Peijie's Definition: Data Science is the science

- *of* the data -- what
- *by* the data -- how
- *for* the data -- why

Mathematics of Data

Representation of Data

What data do we have, and how to relate it with math objects?

Tabular Data

In []:

```
import pandas as pd
import numpy as np
df_house = pd.read_csv('./data/kc_house_data.csv')
print(df_house.shape)
df_house.head()
```

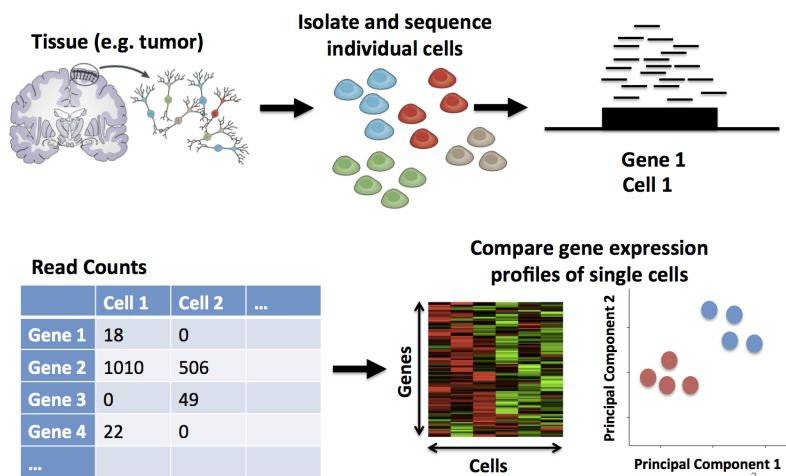
- A structured data table, with n observations and p variables.
- **Mathematical representation:** The data *matrix* $X \in \mathbb{R}^{n \times p}$. For notations we write

$$X = \begin{pmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \dots \\ \mathbf{x}^{(n)} \end{pmatrix}, \text{ where the } i\text{-th row vector represents } i\text{-th observation,}$$

$$\mathbf{x}^{(i)} = (x_1^{(i)}, \dots, x_p^{(i)}) \in \mathbb{R}^p.$$

- [Example: Precision Medicine and Single-cell Sequencing.](#)

Single-cell RNA-Seq (scRNA-Seq)



- Roughly speaking, big data -- large n , high-dimensional data -- large p .

Time-series Data

In []:

```
import matplotlib.pyplot as plt
ts_tesla = pd.read_csv('./data/Tesla.csv')
print(ts_tesla.head())

ts_tesla['Date'] = pd.to_datetime(ts_tesla['Date'])
ts_tesla.set_index('Date', inplace=True)

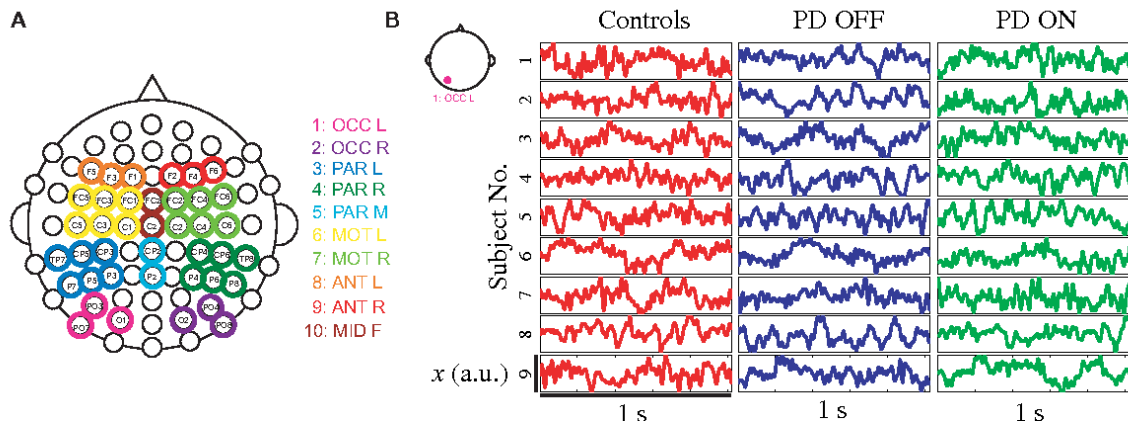
# Suppose we only focus on the time-series of close price
plt.figure(dpi=80)
plt.title('Close Price History')
plt.plot(ts_tesla['Close'], color='red')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD', fontsize = 18)
plt.show()
# this is only about tesla -- we can also have the time-series of apple,amazon,facebook
```

- Simple case: N one-dimensional trajectories with each sampled at T time points.
- **Mathematical representation I:** Still use the data matrix $X \in \mathbb{R}^{N \times T}$. For notations we write

$$X = \begin{pmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \dots \\ \mathbf{x}^{(N)} \end{pmatrix}, \text{ where the } i\text{-th row vector represents } i\text{-th trajectory,}$$

$$\mathbf{x}^{(i)} = (x_1^{(i)}, \dots, x_T^{(i)}) \in \mathbb{R}^T.$$

- Question: The difference with tabular data?
- **Mathematical representation II:** Each trajectory is a *function* of time t . The whole dataset can be represented as $z = f(\omega, t)$ where ω represents the sample and t represents the time. In probability theory, this is called *stochastic process*.
 - For fixed ω , we have a trajectory, which is the function of time.
 - For fixed t , we obtain an ensemble drawn from random distribution.
- Question: How about N d -dimensional trajectories with each sampled at T time points?
- [Example: Electroencephalography \(EEG\) data and Parkinson's disease.](#)



Images

Example: [MNIST handwritten digits data](#): Each image is 28x28 matrix

```
In [ ]: import pandas as pd
mnist = pd.read_csv('./data/train.csv') # stored as data table
mnist.sample(5)
```


```
In [ ]: mnist.shape
```

```
In [ ]: target = mnist['label']
mnist = mnist.drop("label",axis=1)

import matplotlib.pyplot as plt
plt.figure(dpi=100)
for i in range(0,70): #plot the first 70 images
    plt.subplot(7,10,i+1)
    grid_data = mnist.iloc[i,:].to_numpy().reshape(28,28) # reshape from 1d to 2d pixels
    plt.imshow(grid_data,cmap='gray_r', vmin=0, vmax=255)
    plt.xticks([])
    plt.yticks([])
plt.tight_layout()
```

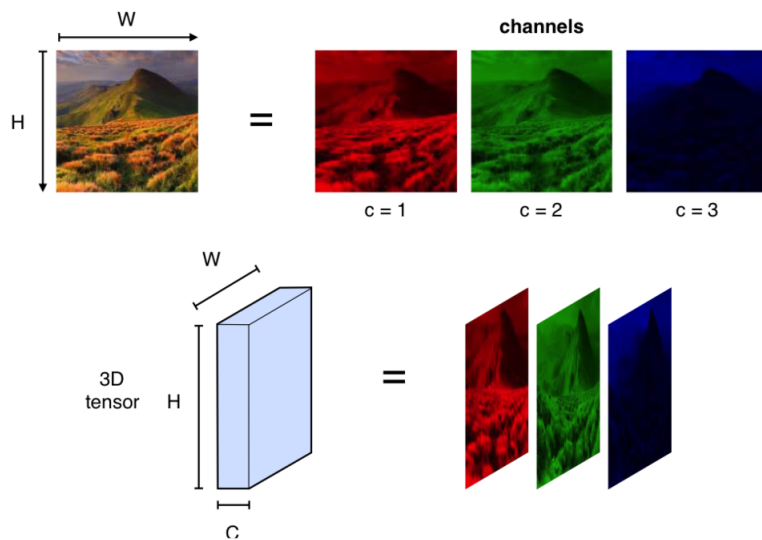
- Simple case: N grayscale images with $m \times n$ pixels each.
- **Mathematical Representation I:** Each image can be represented by a matrix $I \in \mathbb{R}^{m \times n}$, whose elements denotes the intensities of pixels. The whole datasets have N matrices of m by n , or represented by a $N \times m \times n$ tensor.

[Illustrated Introduction to Linear Algebra using NumPy](#)



- **Mathematical representation II:** Random field model $z = \mathbf{f}(\omega, x, y)$.
- **Color images:** Decompose into RGB (red, green and blue) channels and
 - use three matrices (or three-dimensional tensor) to represent one image, or
 - build the random field model with vector-valued functions $z = \mathbf{f}(\omega, x, y) \in \mathbb{R}^3$

[convolutional neural networks](#)



- Question: Can image datasets also be transformed into tabular data? What are the pros/cons?

```
In [ ]: mnist.head()
```

Videos

- *Time-series of images, or random field model $z = \mathbf{f}(\omega, x, y, t)$*

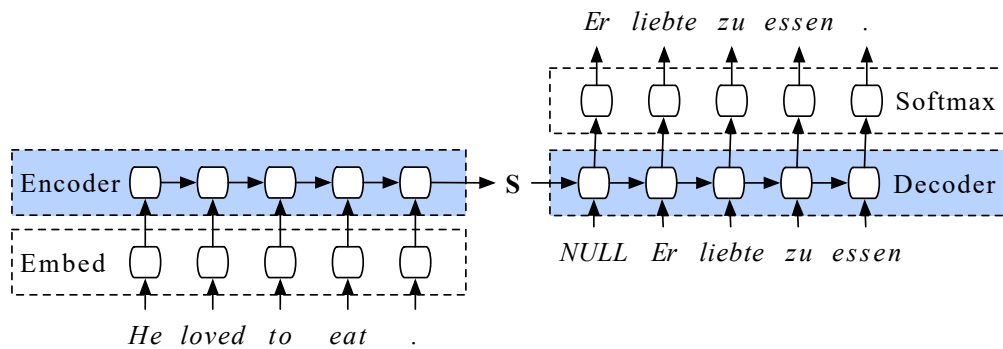
Texts

```
In [ ]: from sklearn.feature_extraction.text import CountVectorizer

corpus = ['He is a good person',
          'He is bad student',
          'He is hardworking']
df = pd.DataFrame(data=corpus, columns=['sentences'])
print(df)
vectorizer = CountVectorizer(vocabulary=['he', 'is', 'a', 'good', 'person', 'bad', 'stu',
                                         stop_words=frozenset(), token_pattern=r"(?u)\b\w+\b")
X = vectorizer.fit_transform(df['sentences'].values)
result = pd.DataFrame(data=X.toarray(), columns=vectorizer.get_feature_names())
result.head()
```

- **Proposal I:** Tabular data by extracting key words. "Document-Term Matrix"
 - useful in sentiment analysis, document clustering, topic modelling
 - popular algorithms include tf-idf, Word2Vec, bag of words, etc.
- **Proposal II:** Time-series of individual words.
 - useful in machine translation

[Recurrent neural network model for machine translations](#)



Networks

- Concepts: node/edge/weight, directed/undirected
- **Mathematical Representation:** adjacency matrix
- Question: what about the whole datasets of networks, and time-evolving networks?

Tasks with Data: Machine Learning

The tasks with data can often be transformed into *machine learning* problems, which can be generally classified as:

- Supervised Learning -- "learning with training";
- Unsupervised Learning -- "learning without training";
- Reinforcement Learning -- "learning by doing".

Our course will focus on the first two categories.

Supervised Learning

- Given the *training dataset* $(x^{(i)}, y^{(i)})$ with $y^{(i)} \in \mathbb{R}^q$ denotes the *labels*, the supervised learning aims to find a mapping $\mathbf{f} : \mathbb{R}^p \rightarrow \mathbb{R}^q$ such that $y^{(i)} \approx \mathbf{f}(x^{(i)})$. Then with a new observation $x^{(new)}$, we can predict that $y^{(new)} = \mathbf{f}(x^{(new)})$.
 - when $y \in \mathbb{R}$ is continuous, the problem is also called as *regression*. **Example:** Housing price prediction
 - when $y \in \mathbb{R}$ is discrete, the problem is also called as *classification*. **Example:** Handwritten digit recognition
- **Practical Strategy:** Limit the mapping \mathbf{f} to certain space by parametrization $\mathbf{f}(\mathbf{x}; \theta)$. Then define the loss function of θ

$$L(\theta) = \sum_{i=1}^n \ell(y^{(i)}, \mathbf{f}(x^{(i)})),$$

where ℓ quantifies the "distance" between $y^{(i)}$ and $\mathbf{f}(x^{(i)})$, and a common choice is mean square error (MSE) for continuous data $\ell(y^{(i)}, \mathbf{f}(x^{(i)})) = \|y^{(i)} - \mathbf{f}(x^{(i)})\|^2$. We then seek to choose the optimal θ that minimizes the loss function

$$\theta^* = \operatorname{argmin}_{\theta} L(\theta),$$

which can be tackled numerically by optimization methods (including the popular stochastic gradient descent).

- Difference choice of $\mathbf{f}(\mathbf{x}; \theta)$ leads to various supervised learning models:
 - Linear function : Linear Regression (for regression)/Logistic Regression (for classification)
 - Composition of linear + nonlinear functions: Neural Network
- **Important Terms:**
 - **Training Data:** Both X and y are provided. The dataset which we use to fit the function.
 - **Test Data:** In principle, only X is provided (some times y^{test} is also provided as the ground-truth to verify). The dataset which we generate new predictions y^{pred} . -- This is the final judgement of your unsupervised ML model!
 - **Validation Data:** A good-fit model on training data does not guarantee the good performance on test data. To gain more confidence before really applying to test data, we "fake" some test data as the "sample exam". To do this, we further split the original training data into new training data and validation data, and then learn the mapping on new training data, and judge on the validation data. We may make some adjustment if the model does not perform well in the "sample exam".
 - Intuitive Understanding: Training data is like quizzes -- you want to learn the "mapping" between the question and correct answer. Test data is like your exam. Validation is like you take a sample exam before the real exam and make some "clinics" about your weakpoints.
 - See the illustration [here](#)

Example: The [Wisconsin breast cancer dataset](#)) and low-code ML package [pycaret](#).

```
In [ ]: pip install --upgrade pycaret #install pycaret -- it's a new package, not coming with A
```

```
In [1]: from sklearn.datasets import load_breast_cancer # Load the dataset
X,y = load_breast_cancer(as_frame = True,return_X_y = True)
```


```
In [3]: X
```

```
Out[3]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dim
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	0.2419	C
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	0.1812	C
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	0.2069	C
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	0.2597	C
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	0.1809	C
...	
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	0.1726	C
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	0.1752	C

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dim
566	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	0.1590	C
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	0.2397	C
568	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	0.1587	C

569 rows × 30 columns



In []:

y

In this dataset, all labels are known. To mimic a real situation, we manually create train and test datasets.

In [4]:

```
from sklearn.model_selection import train_test_split # manually split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=
```

In [5]:

```
X_train.shape
```

Out[5]: (381, 30)

In [6]:

```
y_test.shape
```

Out[6]: (188,)

In [7]:


```
import pandas as pd
data_train = pd.concat([X_train,y_train],axis=1) # the whole data table of training
data_train
```

Out[7]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dim
56	19.210	18.57	125.50	1152.0	0.10530	0.12670	0.13230	0.089940	0.1917	(
144	10.750	14.97	68.26	355.3	0.07793	0.05139	0.02251	0.007875	0.1399	(
60	10.170	14.88	64.55	311.9	0.11340	0.08061	0.01084	0.012900	0.2743	(
6	18.250	19.98	119.60	1040.0	0.09463	0.10900	0.11270	0.074000	0.1794	(
8	13.000	21.82	87.50	519.8	0.12730	0.19320	0.18590	0.093530	0.2350	(
...	
277	18.810	19.98	120.90	1102.0	0.08923	0.05884	0.08020	0.058430	0.1550	(
9	12.460	24.04	83.97	475.9	0.11860	0.23960	0.22730	0.085430	0.2030	(
359	9.436	18.32	59.82	278.6	0.10090	0.05956	0.02710	0.014060	0.1506	(

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dim
192	9.720	18.22	60.73	288.1	0.06950	0.02344	0.00000	0.000000	0.1653	(
559	11.510	23.93	74.52	403.5	0.09261	0.10210	0.11120	0.041050	0.1388	(

381 rows × 31 columns



In [25]:

```
from pycaret.classification import setup
from pycaret.classification import compare_models

bc = setup(data=data_train, target='target') # target is the y column name we want to p
```

	Description	Value
0	session_id	7232
1	Target	target
2	Target Type	Binary
3	Label Encoded	0: 0, 1: 1
4	Original Data	(381, 31)
5	Missing Values	False
6	Numeric Features	30
7	Categorical Features	0
8	Ordinal Features	False
9	High Cardinality Features	False
10	High Cardinality Method	None
11	Transformed Train Set	(266, 29)
12	Transformed Test Set	(115, 29)
13	Shuffle Train-Test	True
14	Stratify Train-Test	False
15	Fold Generator	StratifiedKFold
16	Fold Number	10
17	CPU Jobs	-1
18	Use GPU	False
19	Log Experiment	False
20	Experiment Name	clf-default-name
21	USI	3312
22	Imputation Type	simple

	Description	Value
23	Iterative Imputation Iteration	None
24	Numeric Imputer	mean
25	Iterative Imputation Numeric Model	None
26	Categorical Imputer	constant
27	Iterative Imputation Categorical Model	None
28	Unknown Categoricals Handling	least_frequent
29	Normalize	False
30	Normalize Method	None
31	Transformation	False
32	Transformation Method	None
33	PCA	False
34	PCA Method	None
35	PCA Components	None
36	Ignore Low Variance	False
37	Combine Rare Levels	False
38	Rare Level Threshold	None
39	Numeric Binning	False
40	Remove Outliers	False
41	Outliers Threshold	None
42	Remove Multicollinearity	False
43	Multicollinearity Threshold	None
44	Clustering	False
45	Clustering Iteration	None
46	Polynomial Features	False
47	Polynomial Degree	None
48	Trigonometry Features	False
49	Polynomial Threshold	None
50	Group Features	False
51	Feature Selection	False
52	Features Selection Threshold	None
53	Feature Interaction	False
54	Feature Ratio	False
55	Interaction Threshold	None

	Description	Value
56	Fix Imbalance	False
57	Fix Imbalance Method	SMOTE

```
In [26]: best = compare_models() # pycaret automatically fit different ML models for you, and co
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
et	Extra Trees Classifier	0.9655	0.9887	0.9816	0.9678	0.9733	0.9238	0.9291	0.2650
lda	Linear Discriminant Analysis	0.9581	0.9892	0.9938	0.9482	0.9691	0.9038	0.9121	0.0300
catboost	CatBoost Classifier	0.9581	0.9889	0.9640	0.9722	0.9667	0.9095	0.9137	5.3810
ridge	Ridge Classifier	0.9546	0.0000	0.9938	0.9411	0.9658	0.8979	0.9047	0.0190
qda	Quadratic Discriminant Analysis	0.9546	0.9905	0.9585	0.9702	0.9640	0.9024	0.9035	0.0220
lightgbm	Light Gradient Boosting Machine	0.9543	0.9904	0.9699	0.9617	0.9642	0.9000	0.9052	0.4910
lr	Logistic Regression	0.9511	0.9878	0.9647	0.9601	0.9615	0.8945	0.8977	1.1400
rf	Random Forest Classifier	0.9507	0.9920	0.9640	0.9611	0.9610	0.8933	0.8981	0.2820
ada	Ada Boost Classifier	0.9472	0.9868	0.9522	0.9663	0.9566	0.8884	0.8951	0.1080
xgboost	Extreme Gradient Boosting	0.9432	0.9878	0.9522	0.9605	0.9546	0.8780	0.8832	0.2610
nb	Naive Bayes	0.9393	0.9890	0.9643	0.9451	0.9534	0.8658	0.8714	0.0230
gbc	Gradient Boosting Classifier	0.9356	0.9849	0.9522	0.9517	0.9491	0.8597	0.8688	0.1700
dt	Decision Tree Classifier	0.9093	0.9031	0.9228	0.9396	0.9277	0.8040	0.8143	0.0260
knn	K Neighbors Classifier	0.8984	0.9629	0.9404	0.9074	0.9215	0.7768	0.7862	0.0550
svm	SVM - Linear Kernel	0.8313	0.0000	0.8632	0.8880	0.8567	0.6414	0.6724	0.0250

```
In [27]: best # the best model selected by pycaret
```

```
Out[27]: ExtraTreesClassifier(bootstrap=False, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
                             oob_score=False, random_state=7232, verbose=0,
                             warm_start=False)
```

```
In [28]: from pycaret.classification import predict_model
         predict_model(best); # predict on the validation data that pycaret have selected -- sam
```

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
-------	----------	-----	--------	-------	----	-------	-----

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Extra Trees Classifier	0.9565	0.9855	0.9706	0.9565	0.9635	0.9097	0.9099


```
In [29]: from pycaret.classification import finalize_model
best_final = finalize_model(best) # re-train the dataset with whole input training data
```

```
In [30]: from pycaret.classification import predict_model
predictions = predict_model(best_final, data = X_test) # make new predictions on new-co
predictions
```

```
Out[30]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	1 dime
512	13.40	20.52	88.64	556.7	0.11060	0.14690	0.14450	0.08172	0.2116	0.
457	13.21	25.25	84.10	537.9	0.08791	0.05205	0.02772	0.02068	0.1619	0.
439	14.02	15.66	89.59	606.5	0.07966	0.05581	0.02087	0.02652	0.1589	0.
298	14.26	18.17	91.22	633.1	0.06576	0.05220	0.02475	0.01374	0.1635	0.
37	13.03	18.42	82.61	523.8	0.08983	0.03766	0.02562	0.02923	0.1467	0.
...
100	13.61	24.98	88.05	582.7	0.09488	0.08511	0.08625	0.04489	0.1609	0.
336	12.99	14.23	84.08	514.3	0.09462	0.09965	0.03738	0.02098	0.1652	0.
299	10.51	23.09	66.85	334.2	0.10150	0.06797	0.02495	0.01875	0.1695	0.
347	14.76	14.74	94.87	668.7	0.08875	0.07780	0.04608	0.03528	0.1521	0.
502	12.54	16.32	81.25	476.3	0.11580	0.10850	0.05928	0.03279	0.1943	0.

188 rows × 32 columns



```
In [31]: df_compare = pd.concat([predictions['Label'],y_test],axis = 1) # compare with the groun
df_compare
```

```
Out[31]:
```

	Label	target
512	0	0
457	1	1
439	1	1
298	1	1
37	1	1
...
100	0	0

	Label	target
336	1	1
299	1	1
347	1	1
502	1	1

188 rows × 2 columns

```
In [32]: import numpy as np
np.mean(predictions['Label'].to_numpy() == y_test.to_numpy()) # calculate the percentag
```

Out[32]: 0.9627659574468085

```
In [33]: from pycaret.classification import create_model
lr = create_model('lr') # what if we only want the logistic regression model?
```

	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	0.9630	0.9941	1.0000	0.9444	0.9714	0.9189	0.9220
1	0.9630	1.0000	0.9412	1.0000	0.9697	0.9222	0.9250
2	0.9630	0.9941	0.9412	1.0000	0.9697	0.9222	0.9250
3	0.8889	0.9765	0.8824	0.9375	0.9091	0.7666	0.7689
4	0.9630	0.9824	1.0000	0.9444	0.9714	0.9189	0.9220
5	0.9630	0.9941	0.9412	1.0000	0.9697	0.9222	0.9250
6	0.9231	0.9804	0.9412	0.9412	0.9412	0.8301	0.8301
7	0.9615	0.9935	1.0000	0.9444	0.9714	0.9128	0.9162
8	0.9231	0.9625	1.0000	0.8889	0.9412	0.8312	0.8433
9	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000
Mean	0.9511	0.9878	0.9647	0.9601	0.9615	0.8945	0.8977
SD	0.0294	0.0114	0.0390	0.0361	0.0236	0.0628	0.0619

```
In [34]: predict_model(lr) # validation dataset -- sample exam!
```


	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Logistic Regression	0.9391	0.9862	0.9706	0.9296	0.9496	0.8728	0.8741

Out[34]:

	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dir
0	16.850000	94.209999	666.000000	0.08641	0.06698	0.051920	0.027910	0.1409	
1	15.210000	78.010002	457.899994	0.08673	0.06545	0.019940	0.016920	0.1638	

	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dir
2	18.049999	105.099998	813.000000	0.09721	0.11370	0.094470	0.059430	0.1861	
3	19.620001	91.120003	599.500000	0.10600	0.11330	0.112600	0.064630	0.1669	
4	12.350000	69.139999	363.700012	0.08518	0.04721	0.012360	0.013690	0.1449	
...	
110	18.889999	72.169998	396.000000	0.08713	0.05008	0.023990	0.021730	0.2013	
111	20.250000	102.599998	761.299988	0.10250	0.12040	0.114700	0.064620	0.1935	
112	13.840000	82.709999	530.599976	0.08352	0.03735	0.004559	0.008829	0.1453	
113	19.660000	131.100006	1274.000000	0.08020	0.08564	0.115500	0.077260	0.1928	
114	16.850000	79.190002	481.600006	0.08511	0.03834	0.004473	0.006423	0.1215	

115 rows × 32 columns



```
In [35]: final_lr = finalize_model(lr)
```

```
In [36]: predictions_lr = predict_model(final_lr, data = X_test)
np.mean(predictions_lr['Label'].to_numpy() == y_test.to_numpy())
```

```
Out[36]: 0.9627659574468085
```

```
In [20]: from pycaret.classification import tune_model
tuned_lr = tune_model(lr) # fine-tuning the parameters in logistic regression
```

	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	0.8889	0.9882	0.8235	1.0000	0.9032	0.7756	0.7959
1	0.9630	0.9824	0.9412	1.0000	0.9697	0.9222	0.9250
2	0.9630	1.0000	0.9412	1.0000	0.9697	0.9222	0.9250
3	0.9630	1.0000	0.9412	1.0000	0.9697	0.9222	0.9250
4	0.9259	0.9882	1.0000	0.8947	0.9444	0.8344	0.8460
5	0.9259	0.9432	1.0000	0.8889	0.9412	0.8421	0.8528
6	0.9615	1.0000	1.0000	0.9412	0.9697	0.9172	0.9204
7	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000
8	0.9615	1.0000	1.0000	0.9412	0.9697	0.9172	0.9204
9	0.8846	0.9750	0.8750	0.9333	0.9032	0.7607	0.7632
Mean	0.9437	0.9877	0.9522	0.9599	0.9541	0.8814	0.8874
SD	0.0347	0.0171	0.0586	0.0433	0.0296	0.0715	0.0677


```
In [21]: predict_model(tuned_lr) # still doing the sample exam -- validation dataset
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Logistic Regression	0.9739	0.9968	0.9718	0.9857	0.9787	0.9450	0.9452

Out[21]:

	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	dir
0	33.810001	70.790001	386.799988	0.07780	0.03574	0.004967	0.006434	0.1845	
1	15.690000	92.680000	664.900024	0.07618	0.03515	0.014470	0.018770	0.1632	
2	19.040001	71.800003	394.100006	0.08139	0.04701	0.037090	0.022300	0.1516	
3	12.220000	65.750000	321.600006	0.09996	0.07542	0.019230	0.019680	0.1800	
4	16.850000	79.190002	481.600006	0.08511	0.03834	0.004473	0.006423	0.1215	
...	
110	17.120001	121.400002	1077.000000	0.10540	0.11000	0.145700	0.086650	0.1966	
111	16.950001	96.220001	685.900024	0.09855	0.07885	0.026020	0.037810	0.1780	
112	18.870001	118.699997	1027.000000	0.09746	0.11170	0.113000	0.079500	0.1807	
113	12.840000	74.339996	412.600006	0.08983	0.07525	0.041960	0.033500	0.1620	
114	18.600000	81.089996	481.899994	0.09965	0.10580	0.080050	0.038210	0.1925	

115 rows × 32 columns



```
In [22]: final_tuned_lr = finalize_model(tuned_lr) #retrain with the whole dataset
```

```
In [23]: predictions_tuned_lr = predict_model(final_tuned_lr, data = X_test)
np.mean(predictions_tuned_lr['Label'].to_numpy() == y_test.to_numpy())
```

Out[23]: 0.9521276595744681

Let's recap the workflow above (or about general supervised learning)

- The **minimum requirement** is that we have a training dataset with both X and y (also called labels, targets...). We want to **fit the mapping** between x and y with **training dataset** (the process is indeed called training), and making predictions about the new y given new X in the test dataset.
 - *Remark 1:* The true y in test dataset sometimes can also be known, so that we can know the performance the model immediately. But in general, we won't expect this.
 - *Remark 2:* In our course, just to mimic a real-world situation, sometimes we manually create (split) the train or test data.

- (Optional) We may train multiple models or one model with multiple parameters. How can we compare them and gain more confidence about the final test? Sometimes we further split the training dataset into (real) training dataset and **validation dataset** (imagine it as the sample exam), so that we can get instant feedback because we know the true label in validation dataset.
- (Optional) During training, to be more cautious, sometimes we even make more "quizzes" -- that is called **cross-validation** (will talk about the details in the next lecture)
- (Optional) With 10 "quizzes" (10-fold cross-validation) and "one sample exam" (validation data), for instance, we finally pick up the best candidate model. Before applying to the real test dataset, we don't want to waste any sample. Therefore we **finalize** training by picking up the winner model, while updating it with all the samples (including the validation data) in the training dataset.
- Finally, applying the model to test data -- wait and see!

Of course, as a math course, we are not satisfied with merely calling functions in pycaret. In the rest of lectures this quarter, we are going to dig into details of some algorithms and learn more underlying math -- turn the black box of ML into white (at least gray) one!

Unsupervised Learning

It is still challenging to give a general and rigorous definition for unsupervised learning mathematically. Let's focus on more specific tasks.

- Dimension Reduction

Given $X \in \mathbb{R}^{n \times p}$, finding a mapping function $\mathbf{f} : \mathbb{R}^p \rightarrow \mathbb{R}^q (q \ll p)$ such that the low-dimensional coordinates $z^{(i)} = \mathbf{f}(x^{(i)})$ "preserve the information" about $x^{(i)}$.

- Question: Difference with supervised learning?
- Linear mapping: Principle Component Analysis (PCA)
- Nonlinear mapping: Manifold Learning, Autoencoder

```
In [37]: from sklearn.datasets import load_iris
X,y = load_iris(return_X_y = True) # Note that in the hw this week, it's not allowed to
X
```

```
Out[37]: array([[5.1, 3.5, 1.4, 0.2],
 [4.9, 3. , 1.4, 0.2],
 [4.7, 3.2, 1.3, 0.2],
 [4.6, 3.1, 1.5, 0.2],
 [5. , 3.6, 1.4, 0.2],
 [5.4, 3.9, 1.7, 0.4],
 [4.6, 3.4, 1.4, 0.3],
 [5. , 3.4, 1.5, 0.2],
 [4.4, 2.9, 1.4, 0.2],
 [4.9, 3.1, 1.5, 0.1],
 [5.4, 3.7, 1.5, 0.2],
 [4.8, 3.4, 1.6, 0.2],
 [4.8, 3. , 1.4, 0.1],
```

[4.3, 3. , 1.1, 0.1],
[5.8, 4. , 1.2, 0.2],
[5.7, 4.4, 1.5, 0.4],
[5.4, 3.9, 1.3, 0.4],
[5.1, 3.5, 1.4, 0.3],
[5.7, 3.8, 1.7, 0.3],
[5.1, 3.8, 1.5, 0.3],
[5.4, 3.4, 1.7, 0.2],
[5.1, 3.7, 1.5, 0.4],
[4.6, 3.6, 1. , 0.2],
[5.1, 3.3, 1.7, 0.5],
[4.8, 3.4, 1.9, 0.2],
[5. , 3. , 1.6, 0.2],
[5. , 3.4, 1.6, 0.4],
[5.2, 3.5, 1.5, 0.2],
[5.2, 3.4, 1.4, 0.2],
[4.7, 3.2, 1.6, 0.2],
[4.8, 3.1, 1.6, 0.2],
[5.4, 3.4, 1.5, 0.4],
[5.2, 4.1, 1.5, 0.1],
[5.5, 4.2, 1.4, 0.2],
[4.9, 3.1, 1.5, 0.2],
[5. , 3.2, 1.2, 0.2],
[5.5, 3.5, 1.3, 0.2],
[4.9, 3.6, 1.4, 0.1],
[4.4, 3. , 1.3, 0.2],
[5.1, 3.4, 1.5, 0.2],
[5. , 3.5, 1.3, 0.3],
[4.5, 2.3, 1.3, 0.3],
[4.4, 3.2, 1.3, 0.2],
[5. , 3.5, 1.6, 0.6],
[5.1, 3.8, 1.9, 0.4],
[4.8, 3. , 1.4, 0.3],
[5.1, 3.8, 1.6, 0.2],
[4.6, 3.2, 1.4, 0.2],
[5.3, 3.7, 1.5, 0.2],
[5. , 3.3, 1.4, 0.2],
[7. , 3.2, 4.7, 1.4],
[6.4, 3.2, 4.5, 1.5],
[6.9, 3.1, 4.9, 1.5],
[5.5, 2.3, 4. , 1.3],
[6.5, 2.8, 4.6, 1.5],
[5.7, 2.8, 4.5, 1.3],
[6.3, 3.3, 4.7, 1.6],
[4.9, 2.4, 3.3, 1.],
[6.6, 2.9, 4.6, 1.3],
[5.2, 2.7, 3.9, 1.4],
[5. , 2. , 3.5, 1.],
[5.9, 3. , 4.2, 1.5],
[6. , 2.2, 4. , 1.],
[6.1, 2.9, 4.7, 1.4],
[5.6, 2.9, 3.6, 1.3],
[6.7, 3.1, 4.4, 1.4],
[5.6, 3. , 4.5, 1.5],
[5.8, 2.7, 4.1, 1.],
[6.2, 2.2, 4.5, 1.5],
[5.6, 2.5, 3.9, 1.1],
[5.9, 3.2, 4.8, 1.8],
[6.1, 2.8, 4. , 1.3],
[6.3, 2.5, 4.9, 1.5],
[6.1, 2.8, 4.7, 1.2],
[6.4, 2.9, 4.3, 1.3],
[6.6, 3. , 4.4, 1.4],
[6.8, 2.8, 4.8, 1.4],
[6.7, 3. , 5. , 1.7],

[6. , 2.9, 4.5, 1.5],
[5.7, 2.6, 3.5, 1.],
[5.5, 2.4, 3.8, 1.1],
[5.5, 2.4, 3.7, 1.],
[5.8, 2.7, 3.9, 1.2],
[6. , 2.7, 5.1, 1.6],
[5.4, 3. , 4.5, 1.5],
[6. , 3.4, 4.5, 1.6],
[6.7, 3.1, 4.7, 1.5],
[6.3, 2.3, 4.4, 1.3],
[5.6, 3. , 4.1, 1.3],
[5.5, 2.5, 4. , 1.3],
[5.5, 2.6, 4.4, 1.2],
[6.1, 3. , 4.6, 1.4],
[5.8, 2.6, 4. , 1.2],
[5. , 2.3, 3.3, 1.],
[5.6, 2.7, 4.2, 1.3],
[5.7, 3. , 4.2, 1.2],
[5.7, 2.9, 4.2, 1.3],
[6.2, 2.9, 4.3, 1.3],
[5.1, 2.5, 3. , 1.1],
[5.7, 2.8, 4.1, 1.3],
[6.3, 3.3, 6. , 2.5],
[5.8, 2.7, 5.1, 1.9],
[7.1, 3. , 5.9, 2.1],
[6.3, 2.9, 5.6, 1.8],
[6.5, 3. , 5.8, 2.2],
[7.6, 3. , 6.6, 2.1],
[4.9, 2.5, 4.5, 1.7],
[7.3, 2.9, 6.3, 1.8],
[6.7, 2.5, 5.8, 1.8],
[7.2, 3.6, 6.1, 2.5],
[6.5, 3.2, 5.1, 2.],
[6.4, 2.7, 5.3, 1.9],
[6.8, 3. , 5.5, 2.1],
[5.7, 2.5, 5. , 2.],
[5.8, 2.8, 5.1, 2.4],
[6.4, 3.2, 5.3, 2.3],
[6.5, 3. , 5.5, 1.8],
[7.7, 3.8, 6.7, 2.2],
[7.7, 2.6, 6.9, 2.3],
[6. , 2.2, 5. , 1.5],
[6.9, 3.2, 5.7, 2.3],
[5.6, 2.8, 4.9, 2.],
[7.7, 2.8, 6.7, 2.],
[6.3, 2.7, 4.9, 1.8],
[6.7, 3.3, 5.7, 2.1],
[7.2, 3.2, 6. , 1.8],
[6.2, 2.8, 4.8, 1.8],
[6.1, 3. , 4.9, 1.8],
[6.4, 2.8, 5.6, 2.1],
[7.2, 3. , 5.8, 1.6],
[7.4, 2.8, 6.1, 1.9],
[7.9, 3.8, 6.4, 2.],
[6.4, 2.8, 5.6, 2.2],
[6.3, 2.8, 5.1, 1.5],
[6.1, 2.6, 5.6, 1.4],
[7.7, 3. , 6.1, 2.3],
[6.3, 3.4, 5.6, 2.4],
[6.4, 3.1, 5.5, 1.8],
[6. , 3. , 4.8, 1.8],
[6.9, 3.1, 5.4, 2.1],
[6.7, 3.1, 5.6, 2.4],
[6.9, 3.1, 5.1, 2.3],
[5.8, 2.7, 5.1, 1.9],

```
[6.8, 3.2, 5.9, 2.3],
[6.7, 3.3, 5.7, 2.5],
[6.7, 3. , 5.2, 2.3],
[6.3, 2.5, 5. , 1.9],
[6.5, 3. , 5.2, 2. ],
[6.2, 3.4, 5.4, 2.3],
[5.9, 3. , 5.1, 1.8]]])
```

```
In [38]: from sklearn.decomposition import PCA
pca = PCA(n_components=2) # principle component analysis, reduce 4-dimensional data to 2
X_pca = pca.fit_transform(X)
X_pca
```


```
Out[38]: array([[ -2.68412563,  0.31939725],
 [ -2.71414169, -0.17700123],
 [ -2.88899057, -0.14494943],
 [ -2.74534286, -0.31829898],
 [ -2.72871654,  0.32675451],
 [ -2.28085963,  0.74133045],
 [ -2.82053775, -0.08946138],
 [ -2.62614497,  0.16338496],
 [ -2.88638273, -0.57831175],
 [ -2.6727558 , -0.11377425],
 [ -2.50694709,  0.6450689 ],
 [ -2.61275523,  0.01472994],
 [ -2.78610927, -0.235112 ],
 [ -3.22380374, -0.51139459],
 [ -2.64475039,  1.17876464],
 [ -2.38603903,  1.33806233],
 [ -2.62352788,  0.81067951],
 [ -2.64829671,  0.31184914],
 [ -2.19982032,  0.87283904],
 [ -2.5879864 ,  0.51356031],
 [ -2.31025622,  0.39134594],
 [ -2.54370523,  0.43299606],
 [ -3.21593942,  0.13346807],
 [ -2.30273318,  0.09870885],
 [ -2.35575405, -0.03728186],
 [ -2.50666891, -0.14601688],
 [ -2.46882007,  0.13095149],
 [ -2.56231991,  0.36771886],
 [ -2.63953472,  0.31203998],
 [ -2.63198939, -0.19696122],
 [ -2.58739848, -0.20431849],
 [ -2.4099325 ,  0.41092426],
 [ -2.64886233,  0.81336382],
 [ -2.59873675,  1.09314576],
 [ -2.63692688, -0.12132235],
 [ -2.86624165,  0.06936447],
 [ -2.62523805,  0.59937002],
 [ -2.80068412,  0.26864374],
 [ -2.98050204, -0.48795834],
 [ -2.59000631,  0.22904384],
 [ -2.77010243,  0.26352753],
 [ -2.84936871, -0.94096057],
 [ -2.99740655, -0.34192606],
 [ -2.40561449,  0.18887143],
 [ -2.20948924,  0.43666314],
 [ -2.71445143, -0.2502082 ],
 [ -2.53814826,  0.50377114],
 [ -2.83946217, -0.22794557],
 [ -2.54308575,  0.57941002],
 [ -2.70335978,  0.10770608],
```

[1.28482569, 0.68516047],
[0.93248853, 0.31833364],
[1.46430232, 0.50426282],
[0.18331772, -0.82795901],
[1.08810326, 0.07459068],
[0.64166908, -0.41824687],
[1.09506066, 0.28346827],
[-0.74912267, -1.00489096],
[1.04413183, 0.2283619],
[-0.0087454 , -0.72308191],
[-0.50784088, -1.26597119],
[0.51169856, -0.10398124],
[0.26497651, -0.55003646],
[0.98493451, -0.12481785],
[-0.17392537, -0.25485421],
[0.92786078, 0.46717949],
[0.66028376, -0.35296967],
[0.23610499, -0.33361077],
[0.94473373, -0.54314555],
[0.04522698, -0.58383438],
[1.11628318, -0.08461685],
[0.35788842, -0.06892503],
[1.29818388, -0.32778731],
[0.92172892, -0.18273779],
[0.71485333, 0.14905594],
[0.90017437, 0.32850447],
[1.33202444, 0.24444088],
[1.55780216, 0.26749545],
[0.81329065, -0.1633503],
[-0.30558378, -0.36826219],
[-0.06812649, -0.70517213],
[-0.18962247, -0.68028676],
[0.13642871, -0.31403244],
[1.38002644, -0.42095429],
[0.58800644, -0.48428742],
[0.80685831, 0.19418231],
[1.22069088, 0.40761959],
[0.81509524, -0.37203706],
[0.24595768, -0.2685244],
[0.16641322, -0.68192672],
[0.46480029, -0.67071154],
[0.8908152 , -0.03446444],
[0.23054802, -0.40438585],
[-0.70453176, -1.01224823],
[0.35698149, -0.50491009],
[0.33193448, -0.21265468],
[0.37621565, -0.29321893],
[0.64257601, 0.01773819],
[-0.90646986, -0.75609337],
[0.29900084, -0.34889781],
[2.53119273, -0.00984911],
[1.41523588, -0.57491635],
[2.61667602, 0.34390315],
[1.97153105, -0.1797279],
[2.35000592, -0.04026095],
[3.39703874, 0.55083667],
[0.52123224, -1.19275873],
[2.93258707, 0.3555],
[2.32122882, -0.2438315],
[2.91675097, 0.78279195],
[1.66177415, 0.24222841],
[1.80340195, -0.21563762],
[2.1655918 , 0.21627559],
[1.34616358, -0.77681835],
[1.58592822, -0.53964071],

```
[ 1.90445637,  0.11925069],
[ 1.94968906,  0.04194326],
[ 3.48705536,  1.17573933],
[ 3.79564542,  0.25732297],
[ 1.30079171, -0.76114964],
[ 2.42781791,  0.37819601],
[ 1.19900111, -0.60609153],
[ 3.49992004,  0.4606741 ],
[ 1.38876613, -0.20439933],
[ 2.2754305 ,  0.33499061],
[ 2.61409047,  0.56090136],
[ 1.25850816, -0.17970479],
[ 1.29113206, -0.11666865],
[ 2.12360872, -0.20972948],
[ 2.38800302,  0.4646398 ],
[ 2.84167278,  0.37526917],
[ 3.23067366,  1.37416509],
[ 2.15943764, -0.21727758],
[ 1.44416124, -0.14341341],
[ 1.78129481, -0.49990168],
[ 3.07649993,  0.68808568],
[ 2.14424331,  0.1400642 ],
[ 1.90509815,  0.04930053],
[ 1.16932634, -0.16499026],
[ 2.10761114,  0.37228787],
[ 2.31415471,  0.18365128],
[ 1.9222678 ,  0.40920347],
[ 1.41523588, -0.57491635],
[ 2.56301338,  0.2778626 ],
[ 2.41874618,  0.3047982 ],
[ 1.94410979,  0.1875323 ],
[ 1.52716661, -0.37531698],
[ 1.76434572,  0.07885885],
[ 1.90094161,  0.11662796],
[ 1.39018886, -0.28266094]])
```

In [39]:

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.set() # set the seaborn theme style
figure = plt.figure(dpi=100)
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, s=15, edgecolor='none', alpha=0.5, cmap=plt.cm
plt.xlabel('PC 1')
plt.ylabel('PC 2')
plt.colorbar();
```



- Clustering

Given $X \in \mathbb{R}^{n \times p}$, finding a partition of the dataset into K groups such that

- data within the same group are similar;
- data from different groups are dissimilar.


```
In [40]: from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=0) #call k-means clustering algorithm
y_km = kmeans.fit_predict(X)
y_km # the groups assigned by algorithm
```

```
Out[40]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 0, 0, 0, 0, 2, 0, 0, 0,
0, 0, 0, 2, 2, 0, 0, 0, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 2, 0, 0, 0,
0, 2, 0, 0, 0, 0, 2, 0, 0, 0, 2, 0, 0, 0, 2, 0, 0, 2], dtype=int32)
```

```
In [41]: import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
fig, (ax1, ax2) = plt.subplots(1, 2, dpi=150, figsize=(10,4))

fig1 = ax1.scatter(X_pca[:, 0], X_pca[:, 1], c=y_km, s=15, edgecolor='none', alpha=0.5, c
fig2 = ax2.scatter(X_pca[:, 0], X_pca[:, 1], c=y, s=15, edgecolor='none', alpha=0.5, cmap
ax1.set_title('K-means Clustering')
legend1 = ax1.legend(*fig1.legend_elements(), loc="best", title="Classes")
ax1.add_artist(legend1)
ax2.set_title('True Labels')
legend2 = ax2.legend(*fig2.legend_elements(), loc="best", title="Classes")
ax2.add_artist(legend2)
```

Out[41]: <matplotlib.legend.Legend at 0x7fb5d4ae2810>



Question: What is the difference between clustering and classification? Can you try classification on Iris data with pycaret right now?

```
In [ ]: # try classification with pycaret for Iris data by yourself!
```